

MongoDB – A hands on Intro

Sridhar Nanjundeswaran

Technical Lead, 10Gen

@snanjund

March 12, 2013



This Talk

- Quick introduction to mongoDB
- Hands on mongoDB
 - Using a location-based app as an example
 - https://github.com/sridharn/sfbayazure_2013_03_12
- Deployment to Windows Azure
- 2.4 new features

Why MongoDB? What is it?

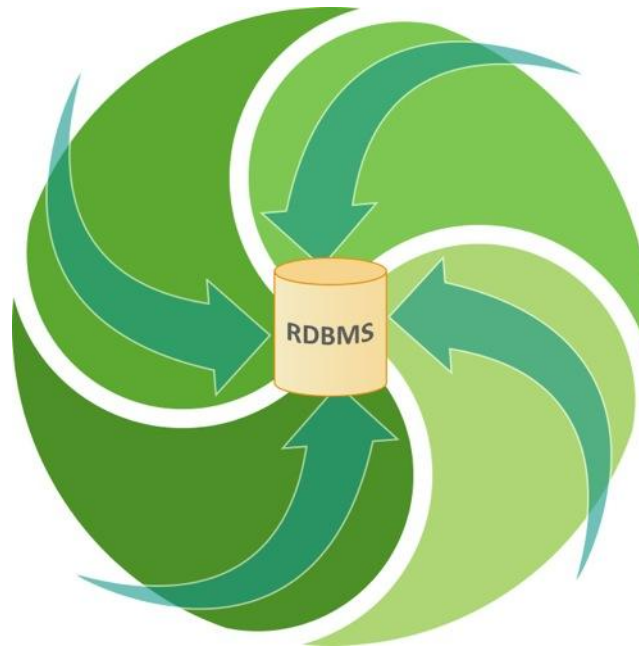
Relational Database Challenges

Data Types

- Unstructured data
- Semi-structured data
- Polymorphic data

Volume of Data

- Petabytes of data
- Trillions of records
- Tens of millions of queries per second



Agile Development

- Iterative
- Short development cycles
- New workloads

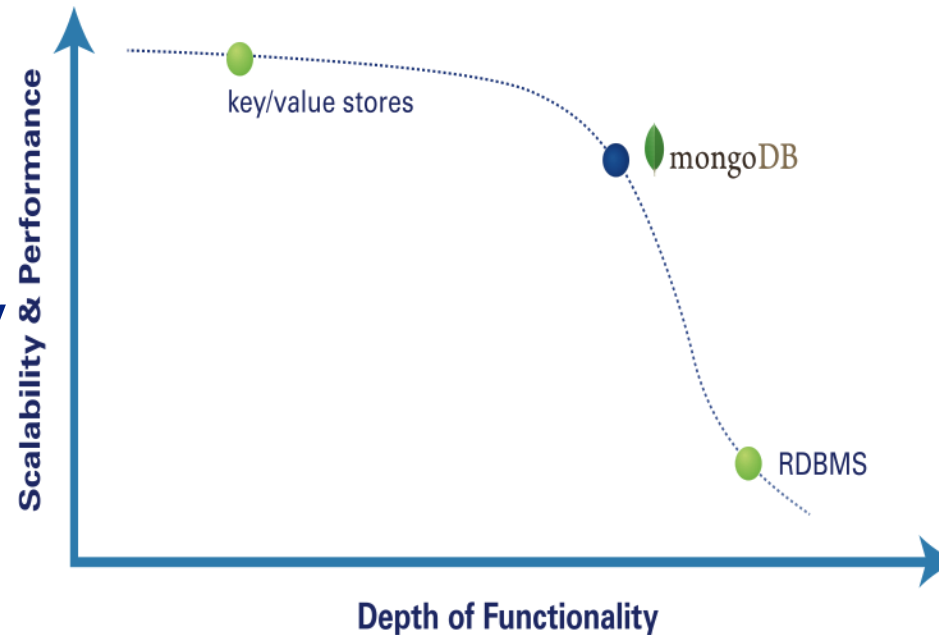
New Architectures

- Horizontal scaling
- Commodity servers
- Cloud computing

What is mongoDB?

MongoDB is a scalable, high-performance, open source, document database.

- Fast Querying
- In-place updates
- Full Index Support
- Replication /High Availability
- Auto-Sharding
- Aggregation; Map/Reduce
- GridFS



BSON

- JSON has powerful, limited set of datatypes
 - Mongo extends datatypes with Date, Int types, ObjectId,
- MongoDB stores data in BSON
- BSON is a binary representation of JSON
 - Optimized for performance and navigational abilities

See: bsonspec.org

BSON { 01010100
11101011
10101110
01010101 }

Where can you use it?

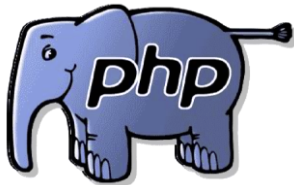
- MongoDB is Implemented in C++
- Windows, Linux, Mac OS-X, Solaris



- Packages available
 - OS X – Macports, Homebrew
 - Linux – Debian, Ubuntu, Fedora, CentOS...

How can I connect to it?

Official 10Gen drivers



MongoDB Drivers

- Official Support for 13 languages
- Community drivers for tons more
 - R, lua etc.
- Drivers connect to mongo servers
- Drivers translate BSON into native types
- mongo shell is not a driver, but works like one in some ways
- Installed using typical means (nuget, npm, pecl, gem, pip)

Terminology

RDBMS	MongoDB
Table	Collection
Row(s)	Document
Index	Index
Partition	Shard
Join	Embedding/Linking
Fixed Schema	Dynamic Schema

Building Your First MongoDB App

- Want to build an app where users can check in to a location



- Leave notes or comments about that location

Requirements

"As a user I want to be able to find other locations nearby"

- Need to store locations (Offices, Restaurants, etc)
 - name, address, tags
 - coordinates
 - User generated content e.g. tips / notes

Requirements

"As a user I want to be able to 'checkin' to a location"

Checkins

- User should be able to 'check in' to a location
- Want to be able to generate statistics:
 - Recent checkins
 - Popular locations

Collections

loc1, loc2, loc3

locations

user1, user2

users

checkin1, checkin2

checkins

Install

- Install DB
- Startup
- Shell
- Code at https://github.com/sridharn/sfbayazure_2013_03_12

Locations v1

```
> location_1 = {  
  name: "Taj Mahal",  
  address: "123 University Ave",  
  city: "Palo Alto",  
  zipcode: 94301  
}
```


Locations v1

```
> location_1 = {  
  name: "Taj Mahal",  
  address: "123 University Ave",  
  city: "Palo Alto",  
  zipcode: 94301  
}  
  
> db.locations.insert(location_1)  
  
> db.locations.find({name: "Taj Mahal"})
```

Locations v1

```
> db.locations.findOne({name: "Taj Mahal"})
```

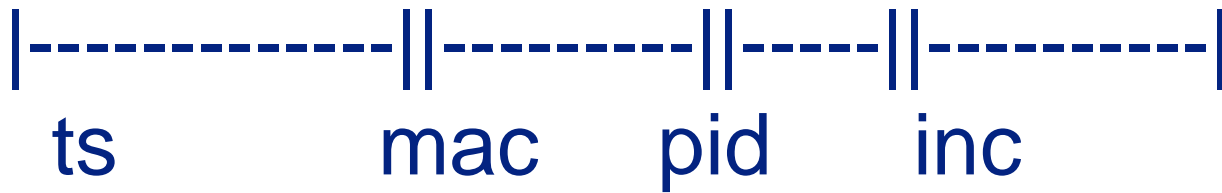
```
{  
  "_id" : ObjectId("50e67a4f4b23019a4ab9b58c"),  
  "name" : "Taj Mahal",  
  "address" : "123 University Ave",  
  "city" : "Palo Alto",  
  "zipcode" : 94301  
}
```

What is _id?

- _id is the primary key in MongoDB
- Automatically indexed
- Automatically created as an ObjectId if not provided
- Any unique immutable value could be used

What is ObjectId?

- ObjectId is a special 12 byte value
- Guaranteed to be unique across your cluster
- ObjectId("50e67a4f4b23019a4ab9b58c")



Locations v1 – Indexed find

```
> db.locations.ensureIndex({name: 1})  
  
> db.locations.find({name: "Taj Mahal"}).explain()  
  
{  
  "cursor" : "BtreeCursor name_1",  
  "isMultiKey" : false,  
  ...
```

Locations v2

```
> location_2 = {  
  name: "Lotus Flower",  
  address: "234 University Ave",  
  city: "Palo Alto",  
  zipcode: 94301,  
  tags: ["restaurant", "dumplings"]  
}
```

```
> db.locations.insert(location_2)
```

Locations v2

```
> db.locations.findOne({tags: "dumplings"})
{
  "_id" : ObjectId("50e67f334b23019a4ab9b59a"),
  "name" : "Lotus Flower",
  ....
}
```

```
> db.locations.ensureIndex({tags: 1})
```

```
> db.locations.find({tags: "dumplings"}).explain()
{
  "cursor" : "BtreeCursor tags_1",
  "isMultiKey" : true,
  ...
}
```

Locations v3

```
> location_3 = {  
  name: "El Capitan",  
  address: "345 University Ave",  
  city: "Palo Alto",  
  zipcode: 94301,  
  tags: ["restaurant", "tacos"],  
  lat_long: [52.5184, 13.387]  
}
```

```
> db.locations.insert(location_3)
```


Locations v3

```
> db.locations.find({lat_long: {$near:[52.53, 13.4]}})
```

```
error: {  
  "$err" : "can't find special index: 2d for: { lat_long: { $near: [ 52.53, 13.4  
] } }",  
  "code" : 13038  
}
```

```
> db.locations.ensureIndex({lat_long: "2d"})
```

```
> db.locations.findOne({lat_long: {$near:[52.53, 13.4]}})  
{  
  "_id" : ObjectId("50e686ab4b23019a4ab9b59d"),  
  "name" : "El Capitan",  
  ...
```

Finding locations

// creating your indexes:

```
> db.locations.ensureIndex({tags: 1})  
> db.locations.ensureIndex({name: 1})  
> db.locations.ensureIndex({lat_long: "2d"})
```

// finding places:

```
> db.locations.find({lat_long: {$near:[52.53, 13.4]}})
```

// with regular expressions:

```
> db.locations.find({name: /^Taj/})
```

// by tag:

```
> db.locations.find({tag: "dumplings"})
```

Updating Documents

Atomic operators:

\$set, \$unset, \$inc, \$push, \$pushAll, \$pull, \$pullAll, \$bit

Inserting locations - adding tips

```
// adding a tip with update:  
> db.locations.update(  
  {name: "Lotus Flower"},  
  {$push: {  
    tips: {  
      user: "Sridhar",  
      date: ISODate("2012-09-21T11:52:27.442Z"),  
      tip: "The sesame dumplings are awesome!"}  
    }  
  })  
})
```

task - done

```
> db.locations.findOne({name:/^Lot/})
{
  "_id" : ObjectId("50e67f334b23019a4ab9b59a"),
  "address" : "234 University Ave",
  "city" : "Palo Alto",
  "name" : "Lotus Flower",
  "tags" : [
    "restaurant",
    "dumplings"
  ],
  "tips" : [
    {
      "user" : "Sridhar",
      "date" : ISODate("2012-09-21T11:52:27.442Z"),
      "tip" : "The sesame dumplings are awesome!"
    }
  ],
  "zipcode" : 94301
}
```



Requirements

"As a user I want to be able to 'checkin' to a location"

Checkins

- User should be able to 'check in' to a location
- Want to be able to generate statistics:
 - Recent checkins
 - Popular locations

Users and Checkins

```
> user_1 = {  
  _id: "sridhar@10gen.com",  
  name: "Sridhar",  
  twitter: "snanjund",  
  checkins: [  
    {location: "Lotus Flower", ts: ISODate("2012-09-21T11:52:27.442Z")},  
    {location: "Taj Mahal", ts: ISODate("2012-09-22T07:15:00.442Z")}  
  ]  
}  
  
> db.users.save(user_1)  
  
> db.users.ensureIndex({"checkins.location": 1})
```

Simple Stats

```
// find all users who've checked in here:  
> db.users.find({"checkins.location":"Lotus Flower"})
```


Simple Stats

// find all users who've checked in here:

```
> db.users.find({"checkins.location":"Lotus Flower"}, {name:1,  
checkins:1})
```

// find the last 10 checkins here?

```
> db.users.find({"checkins.location":"Lotus Flower"}, {name:1,  
checkins:1}).sort({"checkins.ts": -1}).limit(10)
```

Simple Stats

// find all users who've checked in here:

```
> db.users.find({"checkins.location":"Lotus Flower"}, {name:1,  
checkins:1})
```

// find the last 10 checkins here: - **Warning!**

```
> db.users.find({"checkins.location":"Lotus Flower"}, {name:1,  
checkins:1}).sort({"checkins.ts": -1}).limit(10)
```

Hard to query for last 10

User and Checkins v2

```
> user_1 = {  
  _id: "sridhar@10gen.com",  
  name: "Sridhar",  
  twitter: "snanjund",  
}
```

```
> location_id = db.locations.findOne({name:"Taj Mahal"}, {_id:1})["_id"]
```

```
> checkin_1 = {  
  location: location_id,  
  user: "sridhar@10gen.com",  
  ts: ISODate("2012-09-21T11:52:27.442Z")  
}
```

Simple Stats

// find all users who've checked in here:

```
> location_id = db.locations.find({"name": "Lotus Flower"})  
> u_ids = db.checkins.find({location: location_id,  
                           {_id: -1, user: 1}})  
> users = db.users.find({_id: {$in: u_ids}})
```

// find the last 10 checkins here:

```
> db.checkins.find({location: location_id})  
   .sort({ts: -1}).limit(10)
```

// count how many checked in today:

```
> db.checkins.find({location: location_id,  
                  ts: {$gt: midnight}}  
   ).count()
```

Aggregation- in Mongo 2.2

// Find most popular locations

```
> agg = db.checkins.aggregate(  
  {$match: {ts: {$gt: now_minus_3_hrs}}},  
  {$group: {_id: "$location", numEntries: {$sum: 1}}}  
)
```

```
> agg.result  
[{"_id": "Lotus Flower", "numEntries" : 17}]
```

Map Reduce

// Find most popular locations

```
> map_func = function() {  
  emit(this.location, 1);  
}
```

```
> reduce_func = function(key, values) {  
  return Array.sum(values);  
}
```

```
> db.checkins.mapReduce(map_func, reduce_func,  
  {query: {ts: {$gt: now_minus_3_hrs}},  
  out: "result"})
```

```
> db.result.findOne()  
{"_id": "Lotus Flower", "value" : 17}
```

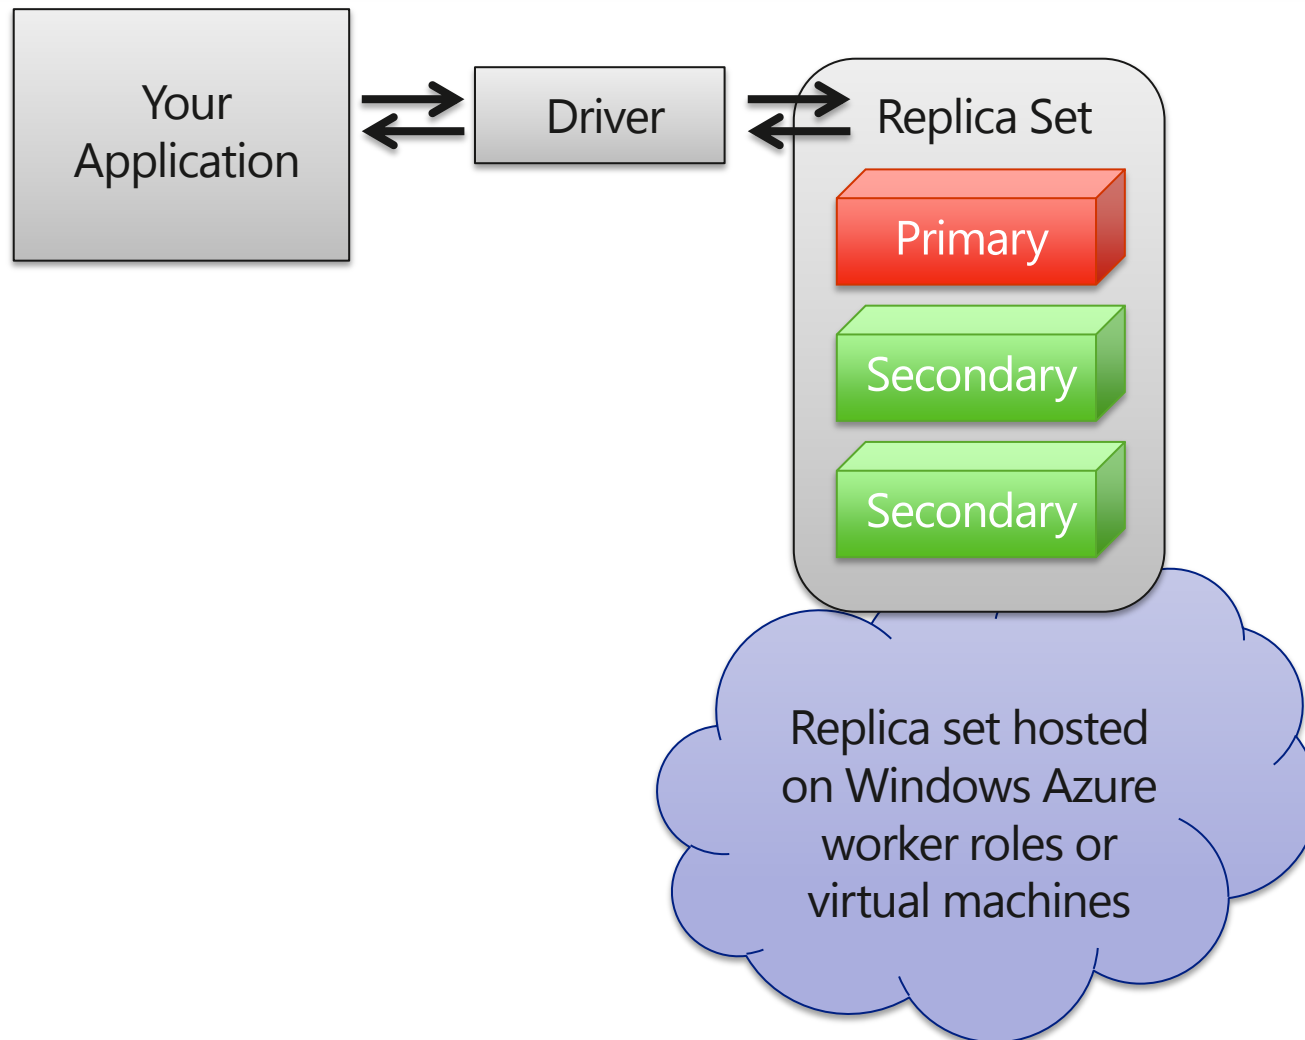
Deployment

MongoDB and Windows Azure

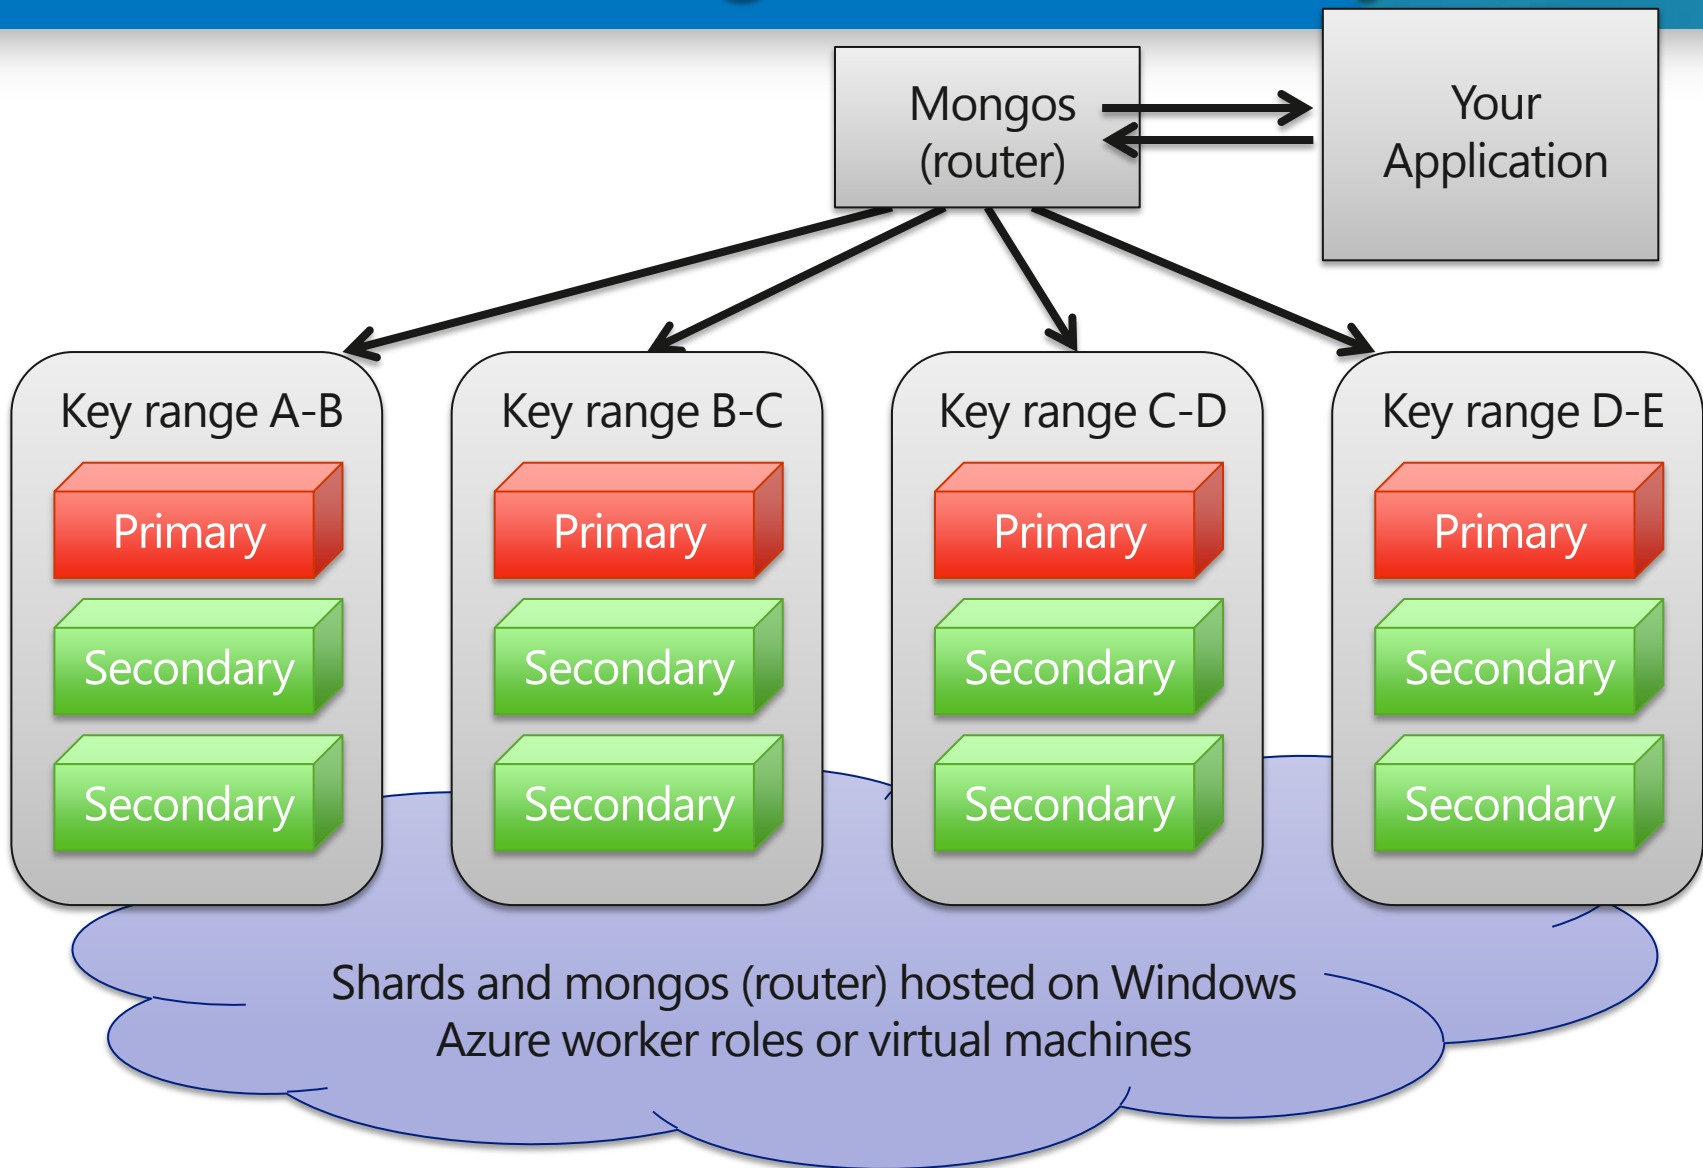
- Choice of popular programming languages
- Deploy on Cloud Services or Virtual Machines

Centers	
mobile	MongoDB drivers:
.net	.NET driver
node.js	Node.js driver
java	Java driver
php	PHP driver
python	Python driver
other	Driver download page



Replica Sets – High availability



Sharding - Scalability



Virtual Machine Sizes

VM Size	CPU Cores	Memory	Bandwidth (Mbps)	# Data Disks	
Extra Small	Shared	768 MB	5	1	
Small	1	1.75 GB	100	2	
Medium	2	3.5 GB	200	4	 
Large	4	7 GB	400	8	
Extra Large	8	14 GB	800	16	

Each Persistent Data Disk Can be up to 1 TB

Windows Azure Cloud Services (PaaS)

- Considerations
 - Instances can reboot
 - IPs can change
 - Input endpoints are load balanced
 - Azure load balancer timeout
- Official Solution
 - Replica set support
 - Should work as is on emulator
 - Only need to configure storage settings
 - Source – <https://www.github.com/mongodb/mongo-azure>
 - Issues – <http://jira.mongodb.org/browse/AZURE>

Demo – Cloud Services (PaaS)

[About](#) [Movies](#)

Index

[Create New](#)

Title	ReleaseDate	Genre	Price	
Lord of the Rings	1/1/2001 12:00:00 AM	Adventure	1.99	Edit Details Delete

App deployed with the MongoDB Replica Set Roles for Windows Azure

Replica Set

Status: OK with replica set name: rs

Id	Name	Health	State	Last Heart Beat	Last Operation	Ping (to Primary)
0	rs_0:27017	● Up	Secondary	9/14/2012 3:49:59 PM	9/14/2012 2:57:20 PM	1
1	rs_1:27017	● Up	👑 Primary	N/A	9/14/2012 2:57:20 PM	N/A
2	rs_2:27017	● Up	Secondary	N/A	9/14/2012 2:57:20 PM	0

Windows Azure Virtual Machines (IaaS)

- Store data in Data Disk
- Multiple instances in a service vs instances across services
- Linux possible
- You need to secure the endpoints
 - Firewall on windows vs linux

Demo – Virtual Machines (IaaS)

```
azure vm create sn0312-c"OpenLogic__OpenLogic-CentOS-62-20120531-  
en-us-30GB.vhd" username password -l "West US" -e
```

```
azure vm create sn0312-c "OpenLogic__OpenLogic-CentOS-62-  
20120531-en-us-30GB.vhd" username password -l "West US" -e 23 -c
```

```
azure vm create sn0312-c "OpenLogic__OpenLogic-CentOS-62-  
20120531-en-us-30GB.vhd" username password -l "West US" -e 24 -c
```

```
azure vm endpoint create sn0312-c 27017 27017
```

```
azure vm endpoint create sn0312-c-2 27018 27018
```

```
azure vm endpoint create sn0312-c-3 27019 27019
```

To set up MongoDB, SSH into each instance and:

- get mongodb binaries and install

- create db dir

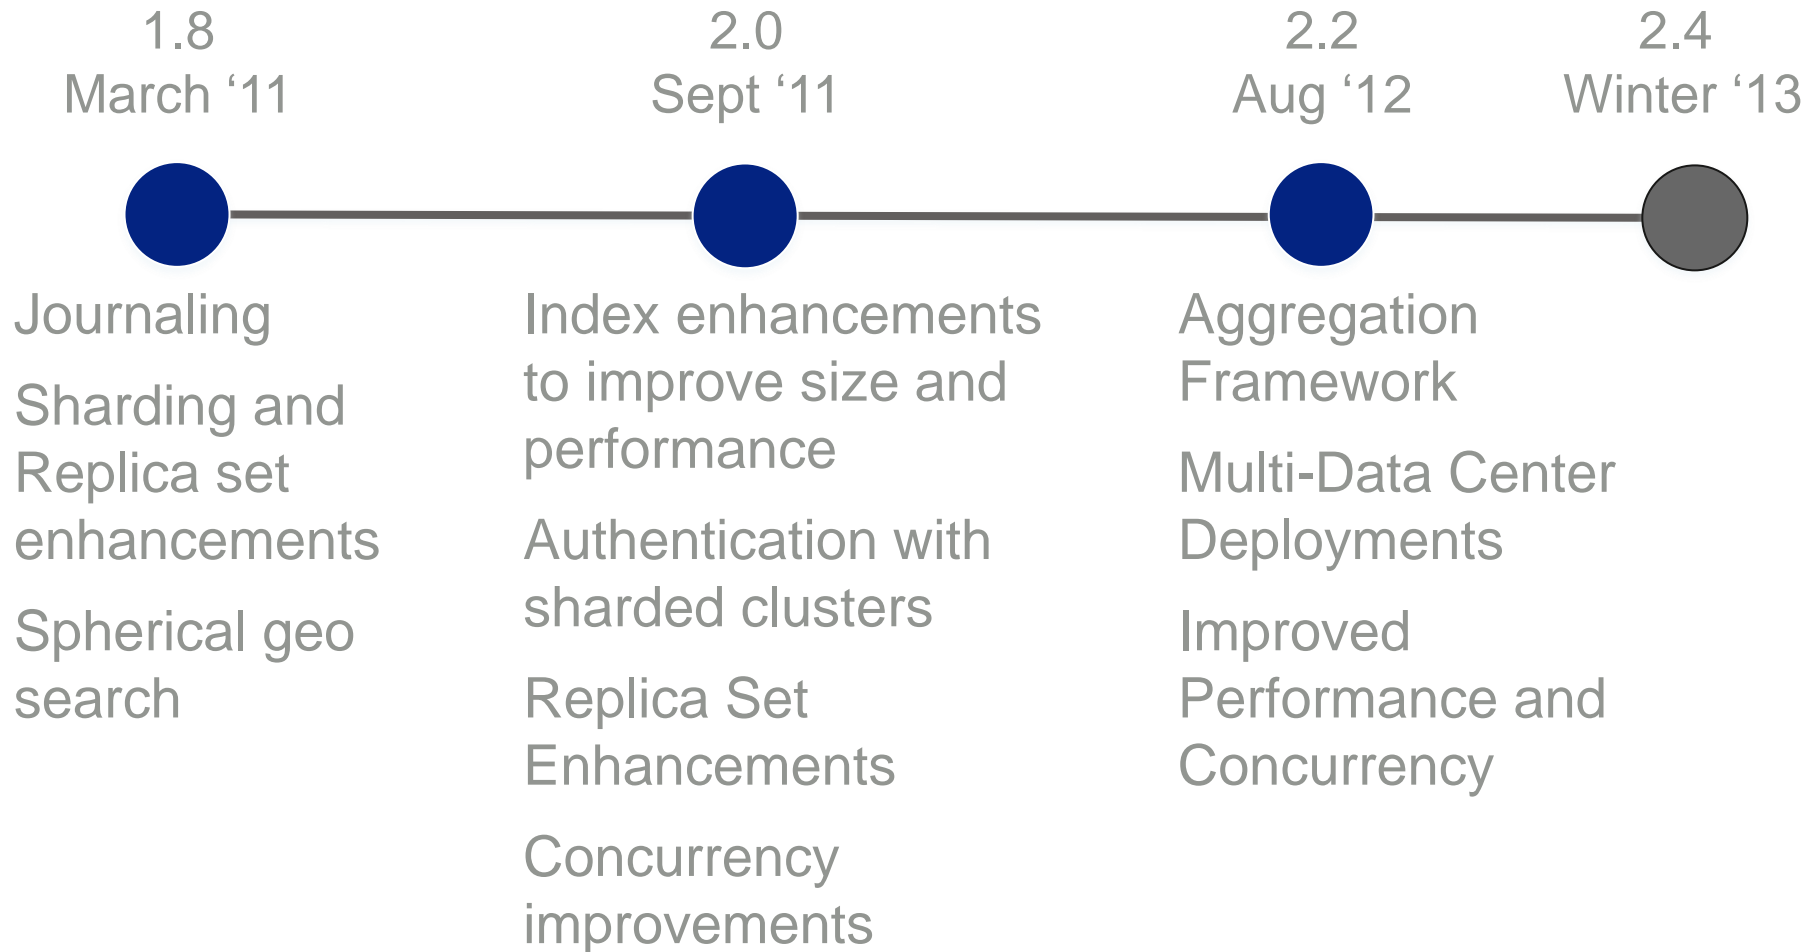
- start mongod, initialize replica set

Resources

- **MongoDB Installer for Windows Azure:**
<http://datamarket.azure.com/application/6ee0c678-21eb-4aff-b345-b371d0c92ecc>
- **MongoDB on Azure:** <http://docs.mongodb.org/ecosystem/platforms/windows-azure/>
- **MongoDB Experts video series:**
<http://blogs.msdn.com/b/interoperability/archive/2012/06/01/mongodb-experts-video-series.aspx>
- **Install MongoDB on a Centos virtual machine in Windows Azure:**
<http://www.windowsazure.com/en-us/manage/linux/common-tasks/mongodb-on-a-linux-vm/>
- **Node.js Web Application with Storage on MongoDB (Virtual Machine):**
[http://waweb.windowsazure.com/en-us/develop/nodejs/tutorials/website-with-mongodb-\(mac\)/](http://waweb.windowsazure.com/en-us/develop/nodejs/tutorials/website-with-mongodb-(mac)/)

Next Steps

Recent Release History



2.2 Release August 2012

- Concurrency: yielding + db level locking
- New aggregation framework
- TTL Collections
- Improved free list implementation
- Tag aware sharding
- Read Preferences
- <http://docs.mongodb.org/manual/release-notes/2.2/>

2.4 Highlights

- Security - SASL, Kerberos, additional privileges
- Hash-based Sharding
- Geospatial Indexing: query intersecting polygons
- Aggregation framework: faster and more features
- V8 JavaScript engine replaces SpiderMonkey
- Replica set flapping reduced, initial sync for replication much faster
- MMS running in your own data center
- Text search (experimental)

Ongoing Work

- Security: finer-level authorization and auditing
- Improved management of sharded clusters
 - Better sharing of network resources in sharded environments (sockets, connections, etc)
- Additional features for Text Search
- Collection-level locking

Where do I go next?

- Questions
 - mongodb-user google group - <https://groups.google.com/forum/?fromgroups#!forum/mongodb-user>
 - Stackoverflow - <http://stackoverflow.com/questions/tagged/mongodb>
 - IRC –freenode.net#mongodb
- Meetups
 - Office hours
 - SF -Epicenter Café alternate Mondays 4 – 6pm
 - Palo Alto – 10Gen office alternate Wednesday 5 – 7pm
 - Meetups
 - SF - <http://www.meetup.com/San-Francisco-MongoDB-User-Group/>
 - Silicon Valley - <http://www.meetup.com/MongoDB-SV-User-Group/>
- Existing bugs and issues
 - <http://jira.mongodb.org>

MongoDB San Francisco

- May 10, 2013 at the Palace Hotel
- Has sold out 3 years in a row
- Still accepting proposals: Submit at 10gen.com/talk-proposal
- Register at mongoSF.com
 - Use discount code `azure10`

Questions?

Sridhar Nanjundeswaran, 10gen
@snanjund

